

Kernels in machine learning


October 2021

Thomas Corpetti




Main principles

- We have separable data, but not linearly separable data




Main principles

- We have separable data, but not linearly separable data
- Make a projection in another space



Main principles

- We have separable data, but not linearly separable data
- Make a projection in another space
- In this new space, data are linearly separable



Outline

1 Introduction

2 What is a kernel ?

- Introduction
- Positive definite kernels
- SVM


3 Evaluation

- Time series

Separability of data


- Ex : find a linear separation of data X in dimension 1

⇒ no solution




Separability of data

- We artificially add a dimension : $X' = (X, \sqrt{X^2})$
⇒ linear solution



Separability of data

- We artificially add a dimension : $X' = (X, \sqrt{X^2})$
 \Rightarrow linear solution



In general, the augmentation on the dimension enables to increase the separability.

Data in large dimension

Make a transformation in a higher dimensional space

- Principle : the initial representation space is transformed :

$$X = [X_1, \dots, X_P] \implies \phi(X) = [\phi(X_1), \dots, \phi(X_P)]$$

with $\phi : \mathbb{R}^N \rightarrow \mathcal{F}^M$ (in general $M \gg N$).

- Difficulty : large computational cost.

Main idea

- Instead of computing the transformation with ϕ on every elements of X , we use a **similarity function** :

$$K : \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$$

$\implies K$ can be seen as a variance-covariance function in the new space (also named *feature space*)

- We try to manipulate only dot products or variance-covariance matrices

K is a kernel function

Data in large dimension

Make a transformation in a higher dimensional space

- Principle : the initial representation space is transformed :

$$X = [X_1, \dots, X_P] \implies \phi(X) = [\phi(X_1), \dots, \phi(X_P)]$$

with $\phi : \mathbb{R}^N \rightarrow \mathcal{F}^M$ (in general $M \gg N$).

- **Difficulty** : large computational cost.

Main idea

- Instead of computing the transformation with ϕ on every elements of X , we use a **similarity function** :

$$K : \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$$

$\implies K$ can be seen as a variance-covariance function in the new space (also named *feature space*)

- We try to manipulate only dot products or variance-covariance matrices

K is a kernel function

Data in large dimension

Make a transformation in a higher dimensional space

- Principle : the initial representation space is transformed :

$$X = [X_1, \dots, X_P] \implies \phi(X) = [\phi(X_1), \dots, \phi(X_P)]$$

with $\phi : \mathbb{R}^N \rightarrow \mathcal{F}^M$ (in general $M \gg N$).

- **Difficulty** : large computational cost.

Main idea

- Instead of computing the transformation with ϕ on every elements of X , we use a **similarity function** :

$$K : \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$$

$\implies K$ can be seen as a variance-covariance function in the new space (also named *feature space*)

- We try to manipulate only dot products or variance-covariance matrices

K is a kernel function

Data in large dimension

Make a transformation in a higher dimensional space

- Principle : the initial representation space is transformed :

$$X = [X_1, \dots, X_P] \implies \phi(X) = [\phi(X_1), \dots, \phi(X_P)]$$

with $\phi : \mathbb{R}^N \rightarrow \mathcal{F}^M$ (in general $M \gg N$).

- **Difficulty** : large computational cost.

Main idea

- Instead of computing the transformation with ϕ on every elements of X , we use a **similarity function** :

$$K : \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$$

$\implies K$ can be seen as a variance-covariance function in the new space (also named *feature space*)

- We try to manipulate only dot products or variance-covariance matrices

K is a kernel function

Data in large dimension

Make a transformation in a higher dimensional space

- Principle : the initial representation space is transformed :

$$X = [X_1, \dots, X_P] \implies \phi(X) = [\phi(X_1), \dots, \phi(X_P)]$$

with $\phi : \mathbb{R}^N \rightarrow \mathcal{F}^M$ (in general $M \gg N$).

- **Difficulty** : large computational cost.

Main idea

- Instead of computing the transformation with ϕ on every elements of X , we use a **similarity function** :

$$K : \mathbb{R}^N \times \mathbb{R}^N \rightarrow \mathbb{R}$$

$\implies K$ can be seen as a variance-covariance function in the new space (also named *feature space*)

- We try to manipulate only dot products or variance-covariance matrices

K is a kernel function

Remarks on kernels

General remarks on kernels

- A **similarity matrix** is a square real matrix, whatever the dimension of data X .
- If a technique involves only similarity matrices, the similarity can be replaced by any other function (**modularity**).
- Whatever the dimension of the transformation ϕ , the kernel always gives a **real function**, with a similarity matrix of size $N \times N$.
- But kernels have to be chosen properly.

Remarks on kernels

General remarks on kernels

- A **similarity matrix** is a square real matrix, whatever the dimension of data X .
- If a technique involves only similarity matrices, the similarity can be replaced by any other function (**modularity**).
- Whatever the dimension of the transformation ϕ , the kernel always gives a **real function**, with a similarity matrix of size $N \times N$.
- But kernels have to be chosen properly.

Remarks on kernels

General remarks on kernels

- A **similarity matrix** is a square real matrix, whatever the dimension of data X .
- If a technique involves only similarity matrices, the similarity can be replaced by any other function (**modularity**).
- Whatever the dimension of the transformation ϕ , the kernel always gives a **real function**, with a similarity matrix of size $N \times N$.
- But kernels have to be chosen properly.

Remarks on kernels

General remarks on kernels

- A **similarity matrix** is a square real matrix, whatever the dimension of data X .
- If a technique involves only similarity matrices, the similarity can be replaced by any other function (**modularity**).
- Whatever the dimension of the transformation ϕ , the kernel always gives a **real function**, with a similarity matrix of size $N \times N$.
- But kernels have to be chosen properly.

Remarks on kernels

Definition

A kernel K defined on \mathbb{R}^N is a **positive-definite kernel** if :

- $K : \mathbb{R}^N \times \mathbb{R}^N \longrightarrow \mathbb{R}$
- $\forall (X, Y) \in (\mathbb{R}^N)^2, K(X, Y) = K(Y, X)$
- $\forall M \in \mathbb{N}, \forall (X_1, \dots, X_M) \in \mathbb{R}^M, \forall (\alpha_1, \dots, \alpha_M) \in \mathbb{R}^M,$

$$\sum_{i=1}^M \sum_{j=1}^M \alpha_i \alpha_j K(X_i, X_j) \geq 0$$

This is equivalent to say that the **similarity matrix** is **positive semi-definite**.

Remarks on kernels

Definition

A kernel K defined on \mathbb{R}^N is a **positive-definite kernel** if :

- $K : \mathbb{R}^N \times \mathbb{R}^N \longrightarrow \mathbb{R}$
- $\forall (X, Y) \in (\mathbb{R}^N)^2, K(X, Y) = K(Y, X)$
- $\forall M \in \mathbb{N}, \forall (X_1, \dots, X_M) \in \mathbb{R}^M, \forall (\alpha_1, \dots, \alpha_M) \in \mathbb{R}^M,$

$$\sum_{i=1}^M \sum_{j=1}^M \alpha_i \alpha_j K(X_i, X_j) \geq 0$$

This is equivalent to say that the **similarity matrix is positive semi-definite**.

Positive definite kernels

Exemples

- Linear kernel :

$$\forall (X, Y) \in (\mathbb{R}^N)^2, K(X, Y) = \langle X, Y \rangle = X^T Y$$

- We have symmetry : $\langle X, Y \rangle = \langle Y, X \rangle$
- We have positivity :

$$\sum_{i=1}^M \sum_{j=1}^M \alpha_i \alpha_j \langle X_i, X_j \rangle = \left\| \sum_{i=1}^M \alpha_i X_i \right\|^2 \geq 0$$

- The projection function is the **identity**

Positive definite kernels

Exemples

- Linear kernel :

$$\forall (X, Y) \in (\mathbb{R}^N)^2, K(X, Y) = \langle X, Y \rangle = X^T Y$$

- We have symmetry : $\langle X, Y \rangle = \langle Y, X \rangle$
- We have positivity :

$$\sum_{i=1}^M \sum_{j=1}^M \alpha_i \alpha_j \langle X_i, X_j \rangle = \left\| \sum_{i=1}^M \alpha_i X_i \right\|^2 \geq 0$$

- The projection function is the **identity**

Positive definite kernels

Exemples

- Polynomial in dimension 1 (if $X, Y \in \mathbb{R}$) :

$$\forall (X, Y) \in \mathbb{R}, K(X, Y) = (1 + XY)^2$$

Positive definite kernels

Examples

- Polynomial in dimension 1 (if $X, Y \in \mathbb{R}$) :

$$\forall (X, Y) \in \mathbb{R}, K(X, Y) = (1 + XY)^2$$

- The projection function is

$$\phi(x) = (1, x\sqrt{2}, x^2)$$

Positive definite kernels



Examples

- Polynomial in dimension 1 (if $X, Y \in \mathbb{R}$) :

$$\forall (X, Y) \in \mathbb{R}, K(X, Y) = (1 + XY)^2$$

- The projection function is

$$\phi(x) = (1, x\sqrt{2}, x^2)$$



Positive definite kernels

Examples

- Polynomial in dimension d (si $X, Y \in \mathbb{R}$) :

$$\forall (X, Y) \in (\mathbb{R})^2, K(X, Y) = (1 + XY)^d$$

The projection function is more tricky to write

Remarks

- If K is a p.s.d, then αK ($\forall \alpha > 0$) is p.s.d too
- If K_1 et K_2 are two p.s.d, then $K_1 + K_2$ is p.s.d too
- If K_1 et K_2 are two p.s.d, then $K_1 K_2$ is p.s.d too

Positive definite kernels

Examples

- Polynomial in dimension d (si $X, Y \in \mathbb{R}$) :

$$\forall (X, Y) \in (\mathbb{R})^2, K(X, Y) = (1 + XY)^d$$

The projection function is more tricky to write

Remarks

- If K is a p.s.d, then αK ($\forall \alpha > 0$) is p.s.d too
- If K_1 et K_2 are two p.s.d, then $K_1 + K_2$ is p.s.d too
- If K_1 et K_2 are two p.s.d, then $K_1 K_2$ is p.s.d too

Translation invariant kernels

Invariance by translation

- Kernels such as $K(X_i, X_j) = Q(X_i - X_j)$
- Classical example : **Gaussian kernel** :

$$K(X_i, X_j) = \exp(-\gamma \|X_i - X_j\|^2)$$

What is the feature space ϕ associated ?

- Many theoretical works on kernel methods. Most common : Gaussian since it projects data in an **infinite dimensional space** (best separability)
- **Kernel trick** : change dot products by kernels

Translation invariant kernels

Invariance by translation

- Kernels such as $K(X_i, X_j) = Q(X_i - X_j)$
- Classical example : **Gaussian kernel** :

$$K(X_i, X_j) = \exp(-\gamma \|X_i - X_j\|^2)$$

What is the feature space ϕ associated ?

- Many theoretical works on kernel methods. Most common : Gaussian since it projects data in an **infinite dimensional space** (best separability)
- **Kernel trick** : change dot products by kernels

Translation invariant kernels

Invariance by translation

- Kernels such as $K(X_i, X_j) = Q(X_i - X_j)$
- Classical example : **Gaussian kernel** :

$$K(X_i, X_j) = \exp(-\gamma \|X_i - X_j\|^2)$$

What is the feature space ϕ associated ?

- Many theoretical works on kernel methods. Most common : Gaussian since it projects data in an **infinite dimensional space** (best separability)
- **Kernel trick** : change dot products by kernels

Translation invariant kernels

Invariance by translation

- Kernels such as $K(X_i, X_j) = Q(X_i - X_j)$
- Classical example : **Gaussian kernel** :

$$K(X_i, X_j) = \exp(-\gamma \|X_i - X_j\|^2)$$

What is the feature space ϕ associated ?

- Many theoretical works on kernel methods. Most common : Gaussian since it projects data in an **infinite dimensional space** (best separability)
- **Kernel trick** : change dot products by kernels

Applications

Many Machine Learning applications use kernels


- Kernel ACP
- Perceptron
- Support Vector Machines
- Kernel Regression
- ...

Applications

Many Machine Learning applications use kernels

- Kernel ACP
- Perceptron
- Support Vector Machines
- Kernel Regression
- ...

Enables to easily deal with non linear data



source : <http://openclassroom.stanford.edu/>

Links between distance and kernels

From a distance function, one can write it with kernels, and conversely

$$\begin{aligned}d(\phi(X), \phi(Y))^2 &= \|\phi(X) - \phi(Y)\|^2 \\&= \phi(X) \cdot \phi(X) - 2\phi(X) \cdot \phi(Y) + \phi(Y) \cdot \phi(Y) \\&= K(X, X) - 2K(X, Y) + K(Y, Y)\end{aligned}$$

Example : distance function for Gaussian kernel ?

$$K(X, Y) = \exp(-\gamma \|X - Y\|^2)$$

Links between distance and kernels

From a distance function, one can write it with kernels, and conversely

$$\begin{aligned}d(\phi(X), \phi(Y))^2 &= \|\phi(X) - \phi(Y)\|^2 \\&= \phi(X) \cdot \phi(X) - 2\phi(X) \cdot \phi(Y) + \phi(Y) \cdot \phi(Y) \\&= K(X, X) - 2K(X, Y) + K(Y, Y)\end{aligned}$$

Example : distance function for Gaussian kernel ?

$$K(X, Y) = \exp(-\gamma \|X - Y\|^2)$$

Outline

1 Introduction

2 What is a kernel ?

- Introduction
- Positive definite kernels
- SVM

3 Evaluation

- Time series

SVM

Main equation

- We aim at **maximizing the margin**, i.e. find the equation of the hyperplane with maximal margin :

$$\frac{1}{2} \min_{\tilde{w}, w_0} \|\tilde{w}\|$$

Under the constraints

$$y_i (\langle \tilde{w}, x_i \rangle + w_0) \geq 1, \quad \forall i \in \{1, \dots, N\}$$

The solution depends only on points called **support vectors**

Important note : only dot products are involved !

SVM

Main equation

- We aim at **maximizing the margin**, i.e. find the equation of the hyperplane with maximal margin :

$$\frac{1}{2} \min_{\tilde{\mathbf{w}}, w_0} \|\tilde{\mathbf{w}}\|$$

Under the constraints

$$y_i (\langle \tilde{\mathbf{w}}, x_i \rangle + w_0) \geq 1, \quad \forall i \in \{1, \dots, N\}$$

The solution depends only on points called **support vectors**

Important note : only dot products are involved !

SVM

Main equation

- We aim at **maximizing the margin**, i.e. find the equation of the hyperplane with maximal margin :

$$\frac{1}{2} \min_{\tilde{\mathbf{w}}, w_0} \|\tilde{\mathbf{w}}\|$$

Under the constraints

$$y_i (\langle \tilde{\mathbf{w}}, x_i \rangle + w_0) \geq 1, \quad \forall i \in \{1, \dots, N\}$$

The solution depends only on points called **support vectors**

Important note : only dot products are involved !

SVM

Main equation

- We aim at **maximizing the margin**, i.e. find the equation of the hyperplane with maximal margin :

$$\frac{1}{2} \min_{\tilde{\mathbf{w}}, w_0} \|\tilde{\mathbf{w}}\|$$

Under the constraints

$$y_i (\langle \tilde{\mathbf{w}}, x_i \rangle + w_0) \geq 1, \quad \forall i \in \{1, \dots, N\}$$

The solution depends only on points called **support vectors**

Important note : only dot products are involved !

SVM

Main equation

- We aim at **maximizing the margin**, i.e. find the equation of the hyperplane with maximal margin :

$$\frac{1}{2} \min_{\tilde{\mathbf{w}}, w_0} \|\tilde{\mathbf{w}}\|$$

Under the constraints


$$y_i (K(\tilde{\mathbf{w}}, x_i) + w_0) \geq 1, \quad \forall i \in \{1, \dots, N\}$$

The solution depends only on points called **support vectors**

Important note : only dot products are involved !


SVM – Illustrations

Illustration (<http://perclass.com/doc/guide/classifiers/svm.html>)




SVM – Illustrations

Illustration (<http://www.ipb.uni-bonn.de/ivm/>)




SVM – Illustrations

Illustration (<http://www.dtrek.com/svm.htm>)



SVM – Illustrations

Illustration (<http://mlpy.sourceforge.net/>)



Support Vector Regression

Principles

- Regression : find a function f that approaches $y \approx f(x)$
- Linear regression : find the vector $\tilde{w} = [w, w_0]^T$ such that
$$f(x) = \langle w, x \rangle + w_0$$

\Rightarrow Least-square
- Linear Support Vector Regression : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{w}\|^2 \\ \text{with constraints } |\langle w, x \rangle + w_0| < \epsilon \end{cases}$$

The solution depends only on support vectors

Support Vector Regression

Principles

- Regression : find a function f that approaches $y \approx f(x)$
- **Linear regression** : find the vector $\tilde{\mathbf{w}} = [\mathbf{w}, w_0]^T$ such that $f(x) = \langle \mathbf{w}, x \rangle + w_0$
 \Rightarrow Least-square
- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{\mathbf{w}}\|^2 \\ \text{with constraints } |\langle \mathbf{w}, x \rangle + w_0| < \epsilon \end{cases}$$

The solution depends only on **support vectors**

Support Vector Regression

Principles

- Regression : find a function f that approaches $y \approx f(x)$
- **Linear regression** : find the vector $\tilde{\mathbf{w}} = [\mathbf{w}, w_0]^T$ such that $f(x) = \langle \mathbf{w}, x \rangle + w_0$
 \Rightarrow Least-square
- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{\mathbf{w}}\|^2 \\ \text{with constraints } |\langle \mathbf{w}, x \rangle + w_0| < \epsilon \end{cases}$$

The solution depends only on **support vectors**

Support Vector Regression

Principles

- Regression : find a function f that approaches $y \approx f(x)$
- **Linear regression** : find the vector $\tilde{\mathbf{w}} = [\mathbf{w}, w_0]^T$ such that $f(x) = \langle \mathbf{w}, x \rangle + w_0$
 \Rightarrow Least-square
- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{\mathbf{w}}\|^2 \\ \text{with constraints } |\langle \mathbf{w}, x \rangle + w_0| < \epsilon \end{cases}$$

The solution depends only on **support vectors**

Support Vector Regression

Principles

- Regression : find a function f that approaches $y \approx f(x)$
- **Linear regression** : find the vector $\tilde{\mathbf{w}} = [\mathbf{w}, w_0]^T$ such that $f(x) = \langle \mathbf{w}, x \rangle + w_0$
 \Rightarrow Least-square
- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{\mathbf{w}}\|^2 \\ \text{with constraints } |\langle \mathbf{w}, x \rangle + w_0| < \epsilon \end{cases}$$

The solution depends only on **support vectors**

Support Vector Regression

Kernel-trick

- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\begin{cases} \min \frac{1}{2} \|\tilde{w}\|^2 \\ \text{with constraints } | \langle \mathbf{w}, x \rangle + w_0 | < \epsilon \end{cases}$$

The solution depends only on **support vectors**

Support Vector Regression

Kernel-trick


- **Linear Support Vector Regression** : similarly than SVM , we look for a hyperplane closed to all data :

$$\left\{ \begin{array}{l} \min \frac{1}{2} \|\tilde{w}\|^2 \\ \text{with constraints } | \langle \mathbf{w}, x \rangle + w_0 | < \epsilon \end{array} \right.$$

The solution depends only on **support vectors**

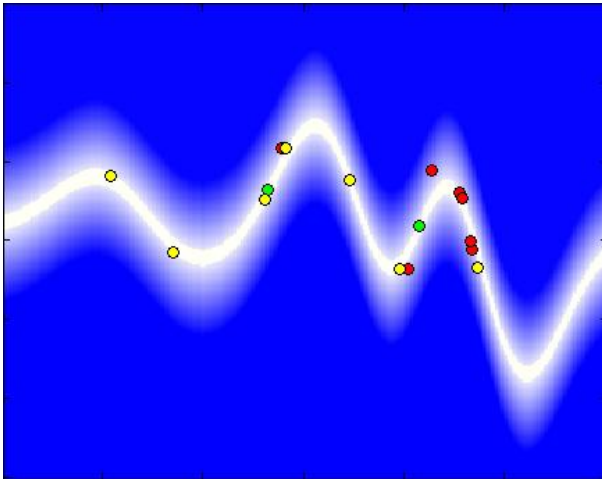
SVR – Illustrations

Illustration (<http://images.1233.tw/support-vector-machine-kernel/>)



SVR – Illustrations

Illustration (<http://onlinesvr.altervista.org/Download.html>)



Evaluation criteria for classification

Precision - Recall - F1-score

- Precision for class i

$$precision(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem assigned to } i}$$

- Overall precision for K classes

$$\frac{\sum_{i=1}^K precision(i)}{K}$$

- Nice if $precision \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - F1-score

- Precision for class i

$$precision(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem assigned to } i}$$

- Overall precision for K classes

$$\frac{\sum_{i=1}^K precision(i)}{K}$$

- Nice if $precision \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - F1-score

- Precision for class i

$$precision(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem assigned to } i}$$

- Overall precision for K classes

$$\frac{\sum_{i=1}^K precision(i)}{K}$$

- Nice if $precision \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - F1-score

- Recall for class i

$$recall(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem in } i}$$

- Overall recall for K classes

$$\frac{\sum_{i=1}^K recall(i)}{K}$$

- Nice if $recall \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - F1-score

- Recall for class i

$$recall(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem in } i}$$

- Overall recall for K classes

$$\frac{\sum_{i=1}^K recall(i)}{K}$$

- Nice if $recall \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - F1-score

- Recall for class i

$$recall(i) = \frac{\text{nb elem correctly assigned to } i}{\text{nb elem in } i}$$

- Overall recall for K classes

$$\frac{\sum_{i=1}^K recall(i)}{K}$$

- Nice if $recall \approx 1$ (but not enough)

Evaluation criteria for classification

Precision - Recall - - F1-score

- Equal combination of recall and precision

$$F1 = 2 \frac{\textit{precision} \cdot \textit{recall}}{\textit{precision} + \textit{recall}}$$

- If precision is good but not recall (we miss data for some classe) : $F1$ decreases (since recall is small)
- If recall is good but not precision (to many data assigned to a class) : $F1$ decreases (since precision is small)
- If recall and precision are good : $F1 \approx 1$

Evaluation criteria for classification

Precision - Recall - - F1-score

- Equal combination of recall and precision

$$F1 = 2 \frac{\textit{precision} \cdot \textit{recall}}{\textit{precision} + \textit{recall}}$$

- If precision is good but not recall (we miss data for some classe) : $F1$ decreases (since recall is small)
- If recall is good but not precision (to many data assigned to a class) : $F1$ decreases (since precision is small)
- If recall and precision are good : $F1 \approx 1$

Evaluation criteria for classification

Precision - Recall - - F1-score

- Equal combination of recall and precision

$$F1 = 2 \frac{\textit{precision} \cdot \textit{recall}}{\textit{precision} + \textit{recall}}$$

- If precision is good but not recall (we miss data for some classe) : $F1$ decreases (since recall is small)
- If recall is good but not precision (to many data assigned to a class) : $F1$ decreases (since precision is small)
- If recall and precision are good : $F1 \approx 1$

Evaluation criteria for classification

Precision - Recall - - F1-score


- Equal combination of recall and precision

$$F1 = 2 \frac{\textit{precision} \cdot \textit{recall}}{\textit{precision} + \textit{recall}}$$

- If precision is good but not recall (we miss data for some classe) : $F1$ decreases (since recall is small)
- If recall is good but not precision (to many data assigned to a class) : $F1$ decreases (since precision is small)
- If recall and precision are good : $F1 \approx 1$

About time series


What possibilities for time series ?



- how to compare them ?

About time series

What possibilities for time series ?



- how to compare them ?

Elastic distances : DTW

Principles : create a path between time series

- Given two series $t_1 = [t_{1,1}, t_{1,2}, \dots, t_{1,m}]^T$ and $t_2 = [t_{2,1}, t_{2,2}, \dots, t_{2,n}]^T$
- Penalization matrix P of size $m \times n$, each element $P(i, j)$ represents the cost to switch from $t_{1,i}$ to $t_{2,j}$:

$$P(i, j) = |t_{1,i} - t_{2,j}|. \quad (1)$$

Example : $t_1 = (2, 3, 6, 9, 5, 4, 3)$ (card = 7) $t_2 = (1, 2, 5, 9, 4, 2)$ (card = 6)

$P(i, j) =$

		t_2						
		1	2	5	9	4	2	$j =$
t_1	2	1	0	3	7	2	0	1
	3	2	1	2	6	1	1	2
	6	5	4	1	3	2	4	3
	9	8	7	4	0	5	7	4
	5	4	3	0	4	1	3	5
	4	3	2	1	5	0	2	6
	3	2	1	2	6	1	1	7
$j =$		1	2	3	4	5	6	

(2)

Elastic distances : DTW

Principles : create a path between time series

- Given two series $t_1 = [t_{1,1}, t_{1,2}, \dots, t_{1,m}]^T$ and $t_2 = [t_{2,1}, t_{2,2}, \dots, t_{2,n}]^T$
- Penalization matrix P of size $m \times n$, each element $P(i, j)$ represents the cost to switch from $t_{1,i}$ to $t_{2,j}$:

$$P(i, j) = |t_{1,i} - t_{2,j}|. \quad (1)$$

Example : $t_1 = (2, 3, 6, 9, 5, 4, 3)$ (card = 7) $t_2 = (1, 2, 5, 9, 4, 2)$ (card = 6)

$P(i, j) =$

		t_2						
		1	2	5	9	4	2	$i =$
t_1	2	1	0	3	7	2	0	1
	3	2	1	2	6	1	1	2
	6	5	4	1	3	2	4	3
	9	8	7	4	0	5	7	4
	5	4	3	0	4	1	3	5
	4	3	2	1	5	0	2	6
	3	2	1	2	6	1	1	7
$j =$		1	2	3	4	5	6	

(2)

Elastic distances : DTW

Principles : create a path between time series

- Given two series $t_1 = [t_{1,1}, t_{1,2}, \dots, t_{1,m}]^T$ and $t_2 = [t_{2,1}, t_{2,2}, \dots, t_{2,n}]^T$
- Penalization matrix P of size $m \times n$, each element $P(i, j)$ represents the cost to switch from $t_{1,i}$ to $t_{2,j}$:

$$P(i, j) = |t_{1,i} - t_{2,j}|. \quad (1)$$

Example : $t_1 = (2, 3, 6, 9, 5, 4, 3)$ (card = 7) $t_2 = (1, 2, 5, 9, 4, 2)$ (card = 6)

$P(i, j) =$

		t_2						
		1	2	5	9	4	2	$i =$
t_1	2	1	0	3	7	2	0	1
	3	2	1	2	6	1	1	2
	6	5	4	1	3	2	4	3
	9	8	7	4	0	5	7	4
	5	4	3	0	4	1	3	5
	4	3	2	1	5	0	2	6
	3	2	1	2	6	1	1	7
$j =$		1	2	3	4	5	6	

(2)

Elastic distances : DTW

Principles : create a path between time series

- A *warping path* has to respect

$W = w_1, \dots, w_K, K \in [\max(m, n), m + n - 1]$:

- $w_1 = (1, 1)$ and $w_K = (m, n)$ (start and end points);
- w_{i+1} is connected w_i for all $i \in [1, K - 1]$ (continuity of the path);
- $(w_{i+1} - w_i)(w_i - w_{i-1}) > 0$ for $i \in [2, K - 1]$ (monotony : no backward path).

Dynamic Time Warping : extract of the minimal cost path to switch from series t_1 to t_2 :

$$D_{dtw}(t_1, t_2) = \min \frac{\sum_{k=1}^K P(w_k)}{K}. \quad (3)$$

Elastic distances : DTW

Principles : create a path between time series

- A *warping path* has to respect

$W = w_1, \dots, w_K, K \in [\max(m, n), m + n - 1]$:

- $w_1 = (1, 1)$ and $w_K = (m, n)$ (start and end points);
- w_{i+1} is connected w_i for all $i \in [1, K - 1]$ (continuity of the path);
- $(w_{i+1} - w_i)(w_i - w_{i-1}) > 0$ for $i \in [2, K - 1]$ (monotony : no backward path).

Dynamic Time Warping : extract of the minimal cost path to switch from series t_1 to t_2 :

$$D_{dtw}(t_1, t_2) = \min \frac{\sum_{k=1}^K P(w_k)}{K}. \quad (3)$$

Elastic distances : DTW

Dynamic Time Warping : extraction of the minimal cost path to switch from series t_1 to t_2 :

$$D_{dtw}(t_1, t_2) = \min \frac{\sum_{k=1}^K P(w_k)}{K}. \quad (4)$$

$P =$

		1	2	5	9	4	2	$i=$
	2	1	0	3	7	2	0	1
	3	2	1	2	6	1	1	2
	6	5	4	1	3	2	4	3
t_1	9	8	7	4	0	5	7	4
	5	4	3	0	4	1	3	5
	4	3	2	1	5	0	2	6
	3	2	1	2	6	1	1	7
	$j=$	1	2	3	4	5	6	


$D =$

		1	2	5	9	4	2	$i=$
	2	1	1	4	11	13	13	1
	3	3	2	3	9	10	11	2
	6	8	6	3	6	8	12	3
t_1	9	16	13	7	3	8	15	4
	5	20	16	7	7	4	7	5
	4	23	18	8	12	4	6	6
	3	25	19	10	14	5	5	7
	$j=$	1	2	3	4	5	6	

(5)

DTW – Illustrations

Illustration (Cassisi et al, 2012)



Average in the DTW distance (useful for k-means)

No analytical formulae

- Average of $[x_1, \dots, x_N]$ w.r.t distance D :

$$\mu = \arg \min_{x^*} \sum_{i=1}^N D^2(x_i, x^*)$$

Average in the DTW distance (useful for k-means)

No analytical formulae

- For DTW :

$$\mu_t = \arg \min_{t^*} \sum_{i=1}^N DTW^2(t_i, t^*)$$


Average in the DTW distance (useful for k-means)

No analytical formulae

- For DTW :

$$\mu_t = \arg \min_{t^*} \sum_{i=1}^N DTW^2(t_i, t^*)$$

Initial curves




Average in the DTW distance (useful for k-means)

No analytical formulae

- For DTW :

$$\mu_t = \arg \min_{t^*} \sum_{i=1}^N DTW^2(t_i, t^*)$$

Possible averages w.r.t. DTW




Average in the DTW distance (useful for k-means)

No analytical formulae

- For DTW :


$$\mu_t = \arg \min_{t^*} \sum_{i=1}^N DTW^2(t_i, t^*)$$

Euclidian mean



Regularized DTW

Illustration DTW : pathological alignments



Regularized DTW

Illustration DTW : pathological alignments

